

# DEEP LEARNING I PROJECT

MASTER OF SCIENCE BIHAR

---

Subject

---

## PREDICTION OF CLOTHING ITEMS USING DEEP LEARNING METHODES (FASHION MNIST DATASET)

---

Presented on February 21, 2022 by

---

**BI TOUVOLY ALAIN SERGE DOUDOU**

**Oumaima ALAMI**

**Hamza BENTAMMAR**

---

Pr Olivier DUBRULE

# SAMMURY

---

1

Logistic Regression

2

Feed-Forward Neural Network

3

Convolutional Neural Network

4

Transfert learning using RESNET16

5

Conclusion

# Logistic Regression

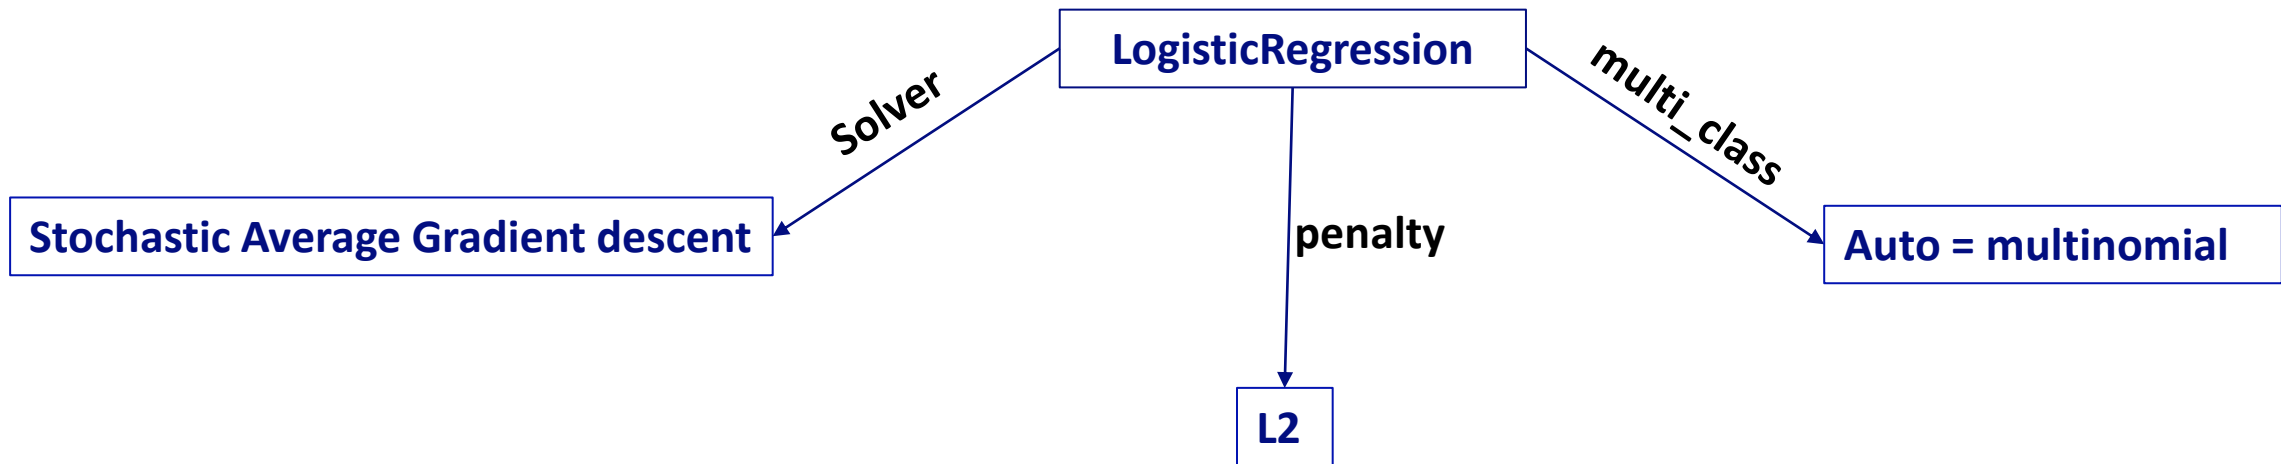
1

# Logistic Regression

Model construction

Model Results

Interpretation




# Logistic Regression

Model construction

Model Results

Interpretation



Model construction

Model Results

Interpretation

	precision	recall	f1-score	support
T-shirt/Top	0.84	0.73	0.78	1000
Trouser	0.84	0.99	0.91	1000
Pullover	0.79	0.60	0.68	1000
Dress	0.93	0.72	0.81	1000
Coat	0.49	0.96	0.65	1000
Sandal	1.00	0.50	0.67	1000
Shirt	0.65	0.29	0.40	1000
Sneaker	0.88	0.73	0.80	1000
Bag	0.83	0.95	0.89	1000
Ankle Boot	0.66	1.00	0.79	1000
accuracy			0.75	10000
macro avg	0.79	0.75	0.74	10000
weighted avg	0.79	0.75	0.74	10000

Logistic Regression achieved an overall accuracy of **75%** on the Fashion-MNIST dataset.

Looking at the classification report, we can see that the model performs relatively well on some classes, such as class 2 achieving an F1-score of 0.91

However, the model struggles with other classes like class 4 (coat) with F1-score 0.65

For further improvement we will try the Feed Forward Neural network approach


# Feed-Forward Neural Network



## Model construction

## Model Results

## Interpretation



Model: "sequential\_1"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense_3 (Dense)	(None, 1024)	803840
dense_4 (Dense)	(None, 10)	10250

=====  
Total params: 814,090  
Trainable params: 814,090  
Non-trainable params: 0  
=====

# Feed-Forward Neural Network

Model construction

Model Results

Interpretation

```
Epoch 1/10
94/94 [=====] - 2s 8ms/step - loss: 0.6000 - accuracy: 0.7967 - val_loss: 0.4692 - val_accuracy: 0.8360
Epoch 2/10
94/94 [=====] - 1s 5ms/step - loss: 0.4140 - accuracy: 0.8571 - val_loss: 0.3956 - val_accuracy: 0.8603
Epoch 3/10
94/94 [=====] - 1s 6ms/step - loss: 0.3690 - accuracy: 0.8706 - val_loss: 0.3739 - val_accuracy: 0.8676
Epoch 4/10
94/94 [=====] - 0s 5ms/step - loss: 0.3377 - accuracy: 0.8788 - val_loss: 0.3531 - val_accuracy: 0.8773
Epoch 5/10
94/94 [=====] - 1s 5ms/step - loss: 0.3210 - accuracy: 0.8840 - val_loss: 0.3425 - val_accuracy: 0.8760
Epoch 6/10
94/94 [=====] - 0s 5ms/step - loss: 0.2996 - accuracy: 0.8912 - val_loss: 0.3313 - val_accuracy: 0.8815
Epoch 7/10
94/94 [=====] - 1s 6ms/step - loss: 0.2876 - accuracy: 0.8938 - val_loss: 0.3222 - val_accuracy: 0.8822
Epoch 8/10
94/94 [=====] - 1s 6ms/step - loss: 0.2733 - accuracy: 0.9002 - val_loss: 0.3092 - val_accuracy: 0.8900
Epoch 9/10
94/94 [=====] - 1s 9ms/step - loss: 0.2620 - accuracy: 0.9040 - val_loss: 0.3107 - val_accuracy: 0.8880
Epoch 10/10
94/94 [=====] - 1s 10ms/step - loss: 0.2511 - accuracy: 0.9097 - val_loss: 0.3127 - val_accuracy: 0.8877
```

Test accuracy: 0.8794999718666077

313/313 [=====] - 1s 2ms/step

Model construction

Model Results

Interpretation


## Classification Report

		precision	recall	f1-score	support
0	T-shirt/Top	0.84	0.83	0.84	1000
1	Trouser	0.99	0.97	0.98	1000
2	Pullover	0.88	0.68	0.77	1000
3	Dress	0.87	0.89	0.88	1000
4	Coat	0.76	0.83	0.79	1000
5	Sandal	0.97	0.97	0.97	1000
6	Shirt	0.65	0.74	0.69	1000
7	Sneaker	0.95	0.95	0.95	1000
8	Bag	0.97	0.97	0.97	1000
9	Ankle Boot	0.96	0.96	0.96	1000
accuracy	accuracy			0.88	10000
macro avg	macro avg	0.88	0.88	0.88	10000
weighted avg	weighted avg	0.88	0.88	0.88	10000

Model construction

Model Results

Interpretation




# Feed-Forward Neural Network

Model construction

Model Results

Interpretation



The feedforward neural network (FFNN) achieved an overall accuracy of **87%** on the Fashion-MNIST dataset.

Looking at the classification report, we can see that the model performs relatively well on some classes, such as class 1 (T-shirt/top) and class 5 achieving an F1-score of 0.99 and 0.93, respectively.

However, the model struggles with other classes, such as class 6 (Shirt), achieving an F1-score of 0.65.

For the validation and training accuracy In the case of the FFNN model trained on the Fashion MNIST dataset, **the training accuracy steadily increases and reaches around 93%** by the end of the 10th epoch, while **the validation accuracy reaches a maximum of around 89%** at around the 7th epoch, and then starts to decrease slightly → This suggests that the model *may be overfitting* the training data, and may not be generalizing well to the validation set or new data.

# Convolutional Neural Network

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
dropout (Dropout)	(None, 14, 14, 32)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 32)	200736
dense_1 (Dense)	(None, 10)	330

=====

Total params: 201,386  
 Trainable params: 201,386  
 Non-trainable params: 0

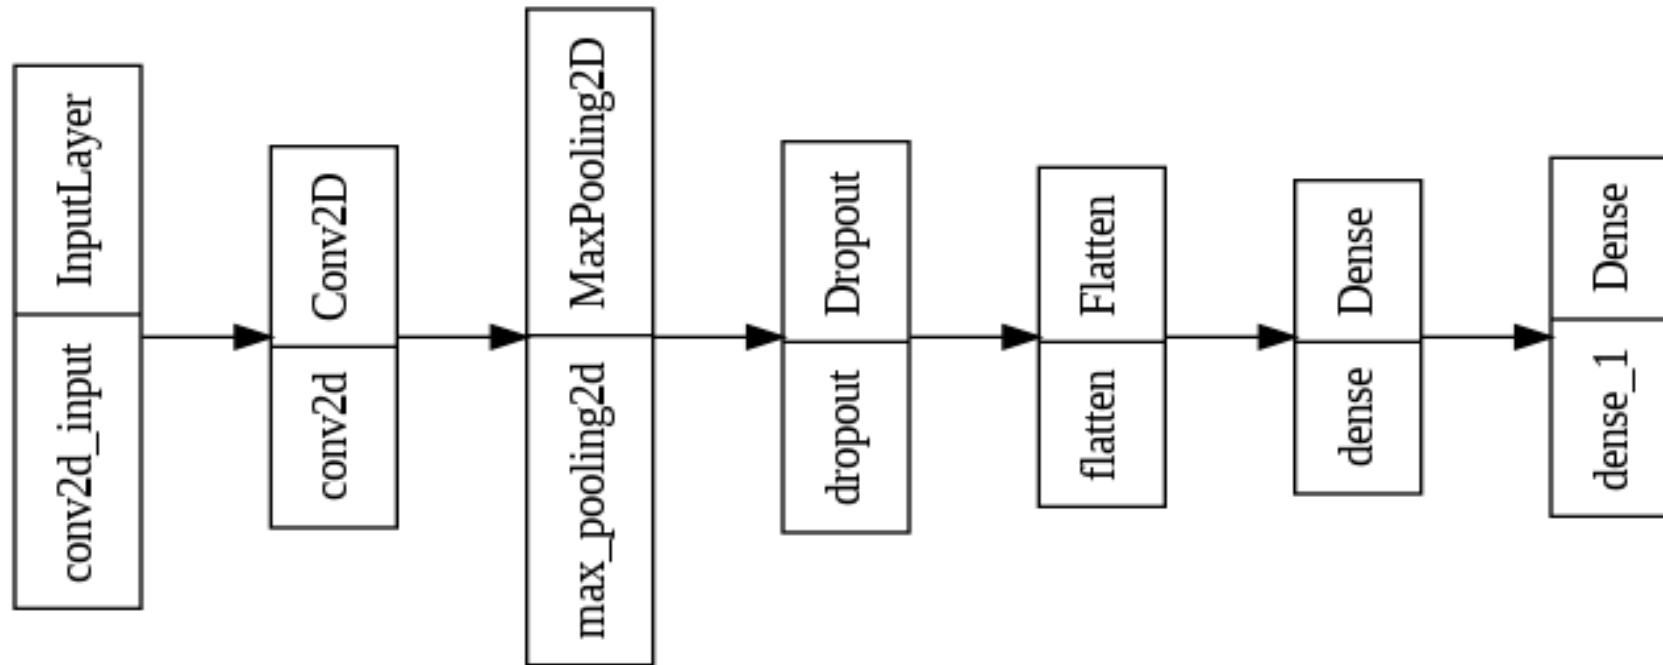


First model

Model construction

Model Results

Interpretation



First model

Model construction

Model Results

Interpretation

```

Epoch 14/20
94/94 [=====] - 32s 341ms/step - loss: 0.2354 - accuracy: 0.9174 - val_loss: 0.2640 - val_accuracy: 0.9030
Epoch 15/20
94/94 [=====] - 33s 347ms/step - loss: 0.2288 - accuracy: 0.9184 - val_loss: 0.2606 - val_accuracy: 0.9044
Epoch 16/20
94/94 [=====] - 32s 341ms/step - loss: 0.2273 - accuracy: 0.9189 - val_loss: 0.2569 - val_accuracy: 0.9060
Epoch 17/20
94/94 [=====] - 35s 368ms/step - loss: 0.2239 - accuracy: 0.9196 - val_loss: 0.2556 - val_accuracy: 0.9072
Epoch 18/20
94/94 [=====] - 31s 331ms/step - loss: 0.2162 - accuracy: 0.9225 - val_loss: 0.2531 - val_accuracy: 0.9068
Epoch 19/20
94/94 [=====] - 33s 352ms/step - loss: 0.2134 - accuracy: 0.9235 - val_loss: 0.2608 - val_accuracy: 0.9062
Epoch 20/20
94/94 [=====] - 32s 343ms/step - loss: 0.2103 - accuracy: 0.9246 - val_loss: 0.2622 - val_accuracy: 0.9047


```

First model

Model construction

Model Results

Interpretation



First model

Model construction

Model Results

Interpretation

	precision	recall	f1-score	support
0	0.83	0.85	0.84	1000
1	0.99	0.96	0.98	1000
2	0.84	0.82	0.83	1000
3	0.86	0.92	0.89	1000
4	0.83	0.82	0.83	1000
5	0.98	0.96	0.97	1000
6	0.70	0.68	0.69	1000
7	0.93	0.95	0.94	1000
8	0.97	0.97	0.97	1000
9	0.95	0.96	0.96	1000
accuracy			0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

**First model****Model construction****Model Results****Interpretation**

The first model of convolutional neural network (CNN) achieved an overall accuracy of **89%** on the Fashion-MNIST dataset.

Looking at the classification report, we can see that the model performs relatively well on some classes, such as class 1 and class 5 achieving an F1-score of 0.98 and 0.97, respectively.

However, the model struggles with other classes, such as class 2, achieving an F1-score of 0.83.

Our model begin to overfit after 14 epochs so we use an other model and see results

Model: "sequential\_1"

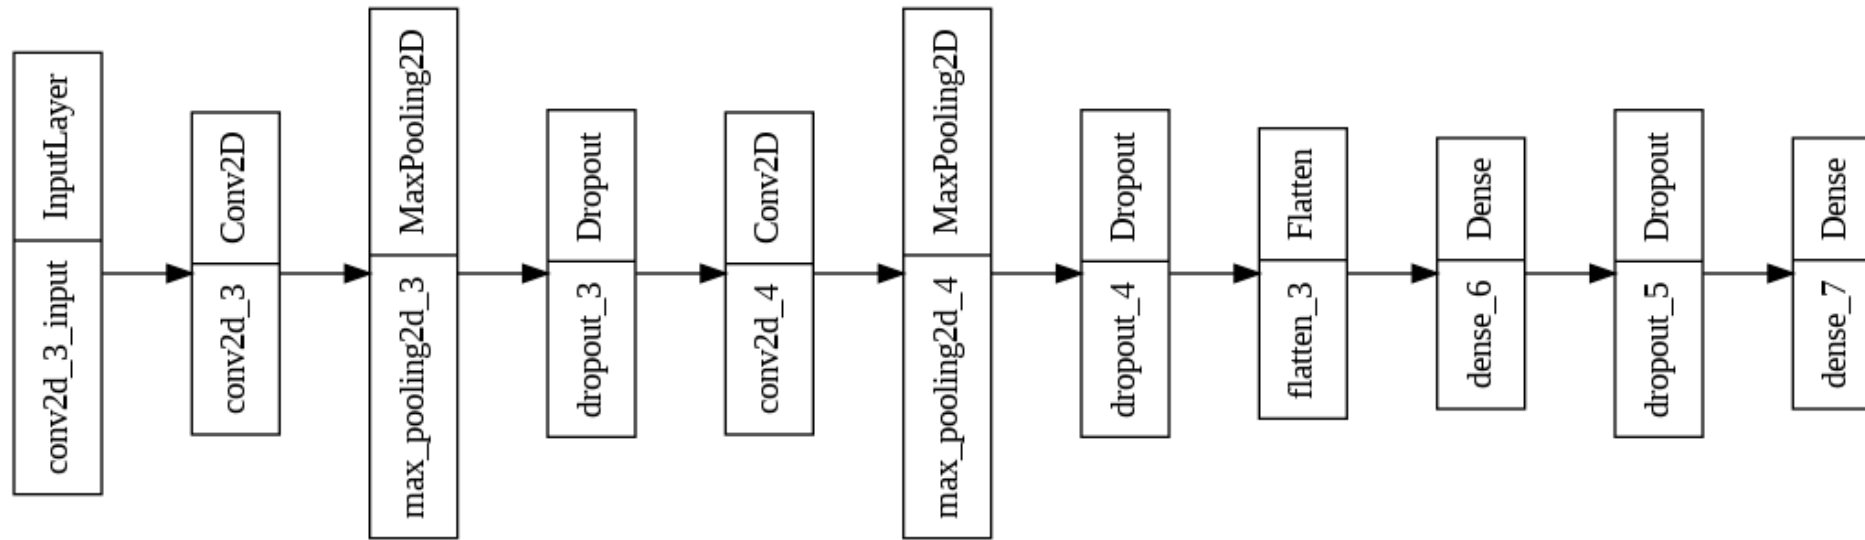
Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_1 (MaxPooling 2D)	(None, 14, 14, 64)	0
dropout_1 (Dropout)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_2 (MaxPooling 2D)	(None, 7, 7, 32)	0
dropout_2 (Dropout)	(None, 7, 7, 32)	0
flatten_1 (Flatten)	(None, 1568)	0
dense_2 (Dense)	(None, 256)	401664
dropout_3 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 10)	2570

=====

Total params: 412,778

Trainable params: 412,778

Non-trainable params: 0




```

Epoch 23/30
94/94 [=====] - 73s 779ms/step - loss: 0.2621 - accuracy: 0.9030 - val_loss: 0.2342 - val_accuracy: 0.9121
Epoch 24/30
94/94 [=====] - 72s 769ms/step - loss: 0.2596 - accuracy: 0.9054 - val_loss: 0.2310 - val_accuracy: 0.9140
Epoch 25/30
94/94 [=====] - 74s 784ms/step - loss: 0.2579 - accuracy: 0.9042 - val_loss: 0.2303 - val_accuracy: 0.9136
Epoch 26/30
94/94 [=====] - 74s 786ms/step - loss: 0.2527 - accuracy: 0.9060 - val_loss: 0.2291 - val_accuracy: 0.9161
Epoch 27/30
94/94 [=====] - 73s 780ms/step - loss: 0.2512 - accuracy: 0.9069 - val_loss: 0.2290 - val_accuracy: 0.9150
Epoch 28/30
94/94 [=====] - 72s 768ms/step - loss: 0.2445 - accuracy: 0.9090 - val_loss: 0.2256 - val_accuracy: 0.9163
Epoch 29/30
94/94 [=====] - 74s 784ms/step - loss: 0.2432 - accuracy: 0.9102 - val_loss: 0.2243 - val_accuracy: 0.9188
Epoch 30/30
94/94 [=====] - 85s 904ms/step - loss: 0.2359 - accuracy: 0.9134 - val_loss: 0.2244 - val_accuracy: 0.9153

```





second model

Model construction

Model Results

Interpretation

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1000
1	0.99	0.98	0.98	1000
2	0.84	0.88	0.86	1000
3	0.91	0.90	0.91	1000
4	0.84	0.84	0.84	1000
5	0.99	0.97	0.98	1000
6	0.76	0.69	0.72	1000
7	0.94	0.98	0.96	1000
8	0.98	0.98	0.98	1000
9	0.97	0.96	0.97	1000
accuracy			0.91	10000
macro avg	0.91	0.91	0.91	10000
weighted avg	0.91	0.91	0.91	10000

The second model of convolutional neural network (CNN) achieved an overall accuracy of **91%** on the Fashion-MNIST dataset.

Looking at the classification report, we can see that the model performs very well on some classes, such as class 1 and class 5 and class 8 achieving an F1-score of 0.98.

However, the model struggles with other classes, such as class 6, achieving an F1-score of 0.72.

The second model gives as a better result but we will try to find the best hyper-parameters.

Best Hyper-Parameters

Best parameters

Model Results

Interpretation

```
optimizer = [SGD, RMSprop, Adagrad, Adadelata, Adam, Adamax, Nadam][:1]

learning_rate = [0.1, 0.001, 0.02][:1]

activation = ['relu', 'tanh', 'sigmoid', 'hard_sigmoid', 'linear'][:1]

num_unit = [10, 5][:1]

initializer = ['lecun_uniform', 'normal', 'he_normal', 'he_uniform'][:1]

rate = [0.3, 0.2, 0.8][:1]

pool_size = [2, 4][:1]

batch_size = [20, 50, 100][:1]

epochs = [10, 20, 50][:1]
```

**Best Hyper-Parameters****Best parameters****Model Results****Interpretation**

```
Best model :  
{'activation': 'relu',  
 'batch_size': 20,  
 'epochs': 10,  
 'initializer': 'lecun_uniform',  
 'learning_rate': 0.1,  
 'num_unit': 10,  
 'optimizer': <class 'keras.optimizers.optimizer_v2.gradient_descent.SGD'>,  
 'pool_size': 2,  
 'rate': 0.3}
```

Best Hyper-Parameters

Best parameters

Model Results

Interpretation

Model: "sequential\_11"

Layer (type)	Output Shape	Param #
conv2d_19 (Conv2D)	(None, 28, 28, 64)	320
max_pooling2d_19 (MaxPooling2D)	(None, 14, 14, 64)	0
dropout_27 (Dropout)	(None, 14, 14, 64)	0
conv2d_20 (Conv2D)	(None, 14, 14, 32)	8224
max_pooling2d_20 (MaxPooling2D)	(None, 7, 7, 32)	0
dropout_28 (Dropout)	(None, 7, 7, 32)	0
flatten_11 (Flatten)	(None, 1568)	0
dense_22 (Dense)	(None, 256)	401664
dropout_29 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 10)	2570

Total params: 412,778

Trainable params: 412,778

Non-trainable params: 0



Best Hyper-Parameters

Best parameters

Model Results

Interpretation

Text(0.5, 1.0, 'Train - Accuracy')



Best Hyper-Parameters

Best parameters

Model Results

Interpretation

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1000
1	0.99	0.98	0.98	1000
2	0.84	0.88	0.86	1000
3	0.91	0.90	0.91	1000
4	0.84	0.84	0.84	1000
5	0.99	0.97	0.98	1000
6	0.76	0.69	0.72	1000
7	0.94	0.98	0.96	1000
8	0.98	0.98	0.98	1000
9	0.97	0.96	0.97	1000
accuracy			0.91	10000
macro avg	0.91	0.91	0.91	10000
weighted avg	0.91	0.91	0.91	10000



Best Hyper-Parameters

Best parameters

Model Results

Interpretation

After running the the model with best parameters, values doesn't change a lot.  
We achieve an overall accuracy of **91%** on the Fashion-MNIST dataset.

# **Transfert learning using RESNET50**

## Model construction

## Model Results

## Interpretation

```

Model: "sequential_12"
Layer (type)                 Output Shape              Param #
=====
resnet50 (Functional)        (None, 2048)              23587712
dense_24 (Dense)              (None, 512)               1049088
dense_25 (Dense)              (None, 10)                 5130
=====
Total params: 24,641,930
Trainable params: 1,054,218
Non-trainable params: 23,587,712

```

# Transfert learning using RESNET50

## Model construction

## Model Results

## Interpretation

```


94/94 [=====] - 5s 58ms/step - loss: 0.1959 - accuracy: 0.9277 - val_loss: 0.4119 - val_accuracy: 0.8593
Epoch 10/20
94/94 [=====] - 6s 59ms/step - loss: 0.1893 - accuracy: 0.9299 - val_loss: 0.4173 - val_accuracy: 0.8574
Epoch 11/20
94/94 [=====] - 6s 59ms/step - loss: 0.1744 - accuracy: 0.9375 - val_loss: 0.4211 - val_accuracy: 0.8642
Epoch 12/20
94/94 [=====] - 5s 58ms/step - loss: 0.1598 - accuracy: 0.9423 - val_loss: 0.4380 - val_accuracy: 0.8617
Epoch 13/20
94/94 [=====] - 5s 58ms/step - loss: 0.1457 - accuracy: 0.9479 - val_loss: 0.4580 - val_accuracy: 0.8587
Epoch 14/20
94/94 [=====] - 5s 55ms/step - loss: 0.1418 - accuracy: 0.9488 - val_loss: 0.4438 - val_accuracy: 0.8657
Epoch 15/20
94/94 [=====] - 5s 58ms/step - loss: 0.1273 - accuracy: 0.9547 - val_loss: 0.4667 - val_accuracy: 0.8621
Epoch 16/20
94/94 [=====] - 5s 57ms/step - loss: 0.1181 - accuracy: 0.9579 - val_loss: 0.4765 - val_accuracy: 0.8639
Epoch 17/20
94/94 [=====] - 5s 58ms/step - loss: 0.1044 - accuracy: 0.9638 - val_loss: 0.4790 - val_accuracy: 0.8586
Epoch 18/20
94/94 [=====] - 5s 54ms/step - loss: 0.1016 - accuracy: 0.9635 - val_loss: 0.4891 - val_accuracy: 0.8640
Epoch 19/20
94/94 [=====] - 6s 59ms/step - loss: 0.0869 - accuracy: 0.9701 - val_loss: 0.4973 - val_accuracy: 0.8652
Epoch 20/20
94/94 [=====] - 5s 55ms/step - loss: 0.0867 - accuracy: 0.9695 - val_loss: 0.5294 - val_accuracy: 0.8621

```

Model construction

Model Results

Interpretation



Model construction

Model Results

Interpretation

- ResNet is a model trained by RGB images and has over 23 millions parameters so it's natural that the model will overfit with our data (grey scale images)
- We will try to do data augmentation in order to help us to deal with this overfitting problem

```
[21] train_generator = ImageDataGenerator(
    rescale = 1./255, # normalization of images
    rotation_range = 40, # augmation of images to avoid overfitting
    shear_range = 0.2,
    zoom_range = 0.2,
    fill_mode = 'nearest')
val_generator = ImageDataGenerator(rescale = 1./255)

train_iterator = train_generator.flow(x_train, y_train, batch_size=512, shuffle=True)

val_iterator = val_generator.flow(x_test, y_test, batch_size=512, shuffle=False)
```

Data augmentation

Model construction

Model Results

Interpretation

```

Epoch 11/20
118/118 [=====] - 29s 249ms/step - loss: 0.8626 - accuracy: 0.6845 - val_loss: 0.7825 - val_accuracy: 0.7100
Epoch 12/20
118/118 [=====] - 30s 257ms/step - loss: 0.8591 - accuracy: 0.6851 - val_loss: 0.7635 - val_accuracy: 0.7178
Epoch 13/20
118/118 [=====] - 29s 250ms/step - loss: 0.8469 - accuracy: 0.6896 - val_loss: 0.7661 - val_accuracy: 0.7116
Epoch 14/20
118/118 [=====] - 30s 252ms/step - loss: 0.8405 - accuracy: 0.6898 - val_loss: 0.7765 - val_accuracy: 0.7062
Epoch 15/20
118/118 [=====] - 30s 253ms/step - loss: 0.8293 - accuracy: 0.6940 - val_loss: 0.7558 - val_accuracy: 0.7145
Epoch 16/20
118/118 [=====] - 30s 250ms/step - loss: 0.8238 - accuracy: 0.6975 - val_loss: 0.7674 - val_accuracy: 0.7114
Epoch 17/20
118/118 [=====] - 29s 250ms/step - loss: 0.8176 - accuracy: 0.6989 - val_loss: 0.7662 - val_accuracy: 0.7066
Epoch 18/20
118/118 [=====] - 29s 249ms/step - loss: 0.8092 - accuracy: 0.7023 - val_loss: 0.7640 - val_accuracy: 0.7133
Epoch 19/20
118/118 [=====] - 29s 247ms/step - loss: 0.8081 - accuracy: 0.7025 - val_loss: 0.7829 - val_accuracy: 0.6918
Epoch 20/20
118/118 [=====] - 30s 255ms/step - loss: 0.8022 - accuracy: 0.7049 - val_loss: 0.7619 - val_accuracy: 0.7065

```




Data augmentation

Model construction

Model Results

Interpretation



Data augmentation

Model construction

Model Results

Interpretation

- After augmenting data using 'DataGenerator' the system doesn't overfit
- but results didn't get any better

Accuracy achieved for Logistic Regression	Accuracy achieved for FFNN	Accuracy achieved for CNN	Accuracy achieved for ResNet50
75%	87%	91%	71%

- We noted that the accuracy achieved by the Logistic regression is not as high the one achieved by FFNN.
- We noted that the accuracy achieved by FFNN model on this dataset is not as high as the CNN model.
- With the ResNet50, we weren't able to find better results due to the high number of parameters and due to the fact that the model is mostly used for the RGB images.

→ The CNN model was the most performant model out of the four